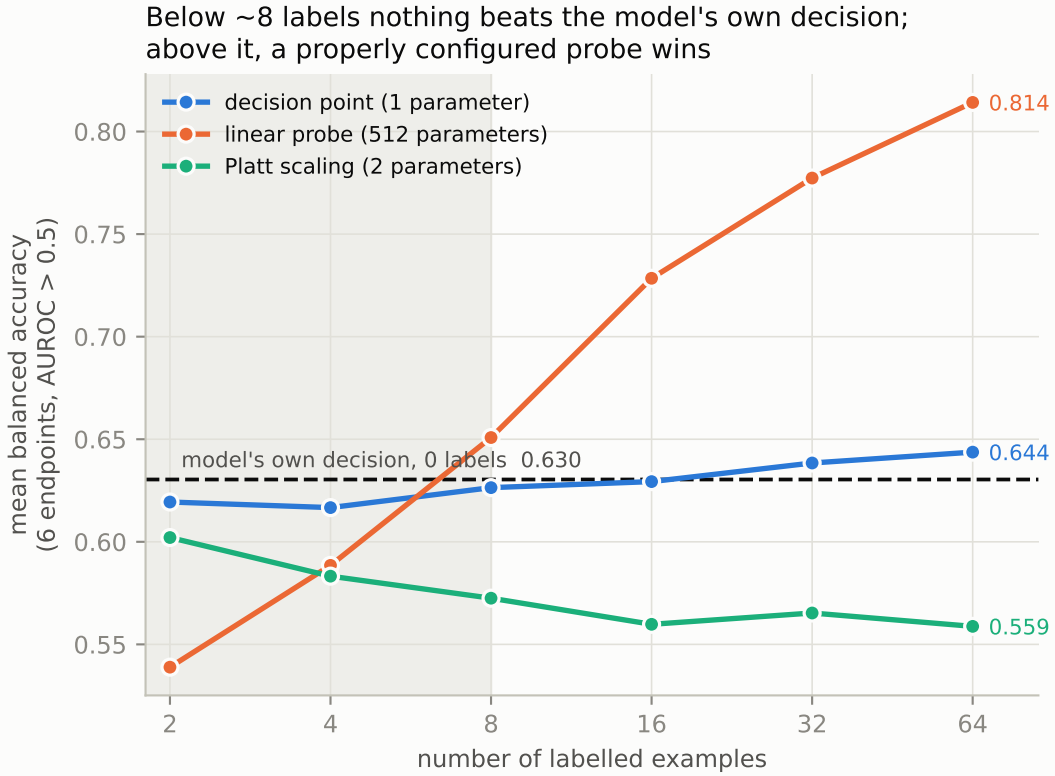
Below ~8 labels nothing beats the model's own decision;
above it, a properly configured probe wins
decision point (1 parameter)
linear probe (512 parameters)
Platt scaling (2 parameters)
0.814
0.644
0.559
model's own decision, 0 labels 0.630
mean balanced accuracy
(6 endpoints, AUROC > 0.5)
0.80
0.75
0.70
0.65
0.60
0.55
2
4
8
16
32
64
number of labelled examples